

Elastic Search

Till now we have looked at:

1. a system that needs sharding vs a system that doesn't need sharding
2. different types of schema
3. Search typeahead : requires hashing
4. Messaging Systems: Column based

How to do a text search?

1. MySQL: indexes keywords
 - a. Resume: Name, id, text
 - b. Find people who have skills: Java / React JS
 - c. like query will work but will not be very optimal
2. NoSQL
 - a. **Key-Value**: Search in the value will be not optimal
 - b. **Document**: Search will not be optimal if we want to query on a part document
 - c. **Column Family**: Same as above, Search will not be optimal if we want to query on a part document
 - d. **Graph**: How will we structure the node and edges?

Google Search

There are three parts to the architecture:

1. Crawler
 - a. Goes to all web pages
 - b. Collects lots of documents
 - c. Billions of web pages
2. Index
 - a. The problem is to store them in such a way, that it is quick to find words matching the search query.
 - b. Search Query: Virat Kohli,
 - c. find all documents where these words appear and give higher priority to pages where they appear together,
 - d. run proprietary algorithm to rank these search result
 - e. For the session we'll work with the subset of the problem:
 - given a phrase find all documents that completely or partially match the query
 - given a single word, how do we find matching pages?
 - given multiple words, how do we find matching pages?

- f. If I search for "quick" and the web page has "quickly", should it match?
 - Yes, but low priority
 - Neglect stopwords: "A" quick fox
 - Number of times the search phrase appear
 - g. substring might not always yield correct results
 - think, thought
 - h. We need some form of mapping
 - word -> document, freq, positions
 - stemming: every words has a root word
 1. run: running, run, ran, bhagna (hindi), odu (Kannada), correr (Spanish)
 2. for gibberish we can't have a root word
 - get rid of insignificant words: articles, prepositions
 1. Ram is running
 2. Ram ran
 - i. Elastic Search is build on Apache lucene, it uses stemming
 - j. Alexa: Voice to text -> text -> semantic -> action (search, buy, volume)
 - k. How should we index?
 - Word -> documents -> frequency per doc + positions of occurrence
 1. Abhishek ate maggi -> Abhishek eat maggi
 2. Higher preference if they occur together
 3. Lower preference if they are scattered in the doc
 - l. Word (Key) -> [{doc_id, position}]
 - run -> [(1,10), (1,20), (2,3)]
 - The array can be quite large, hence, we'll talk about sharding as well
 - This is also called inverted index
3. Query
- a. Search term: Abdul Kalam
 - b. Ranking takes diff params: domain score, number of backlinks Words in H1, H2 etc
 - c. Based on the query we can have three scenarios:
 - d. Docs that contain both Abdul and Kalam: R1
 - e. Docs that contain only Abdul: R2
 - f. Docs that contain only Kalam: R2
 - g. R1 can be split into
 - words appear consecutively: R11
 - words don't appear consecutively: R12
 - h. Sorting: R11 (higher count) > R11 (lower count) > R12 > R2
 - i. Ranking
 - AND (all keywords)

- OR
- j. Tech Design?
 - Abdul -> [(1,10), (1,25), (1,30), (2,15),(2,38), (3,50)]
 - Kalam -> [(1,15) (3,51)]
 - Now we need to find documents that have all keywords
 - how to find documents with occurrences in the consecutive order? // DSA Problem
 1. A1: [(1,10), (1,25), (1,30), (2,15),(2,38), (3,50)]
 2. A2: [(1,15), (3,51)]
 3. variables: i, j , consecutive_cnt[doc_id], non_consecutive_cnt[doc_id]
 4. increment counters and update relevant counters
- k. In real word we have lots of documents
 - array corresponding to a value can be quite large
 - if the size is quite large, how would we shard?
 - word / document_id
 - Two types of queries
 1. Write*: (word, doc)
 2. Read*: word -> fetch matching docs // trying to optimise for this to keep latency low
- l. Availability vs Consistency: Availability

Sharding

Sharding on words:

- Abdul will be on S1
- Kalam will be on S2
- Array can become quite large // 1 billion entries
- Copy all entries to a central place to figure out the common docs // time consuming
- Go to S1, S2 and fetch A1 and A2 // can be done in parallel
- 2 pointer to find intersection // for billion entries, this will also take time
- Update will be slow, since we'll have to update shards corresponding to each word

Sharding on document

- Let's say we have 1000 shards, then every shard has 10 million document
- Every shard has its own independent inverted index
- Every single word from a page will be present on a single shard
- Shard1 can have max 10M entries
- In every shard (in parallel) use 2 pointer approach to find intersection

- Shards reply back with document_id, score
- with whatever response I have in 0.5 seconds, we sort & return the response
- Downside
 - every query will run on every shard
 - cost intensive
 - TPS on each shard will be very high

Extras

Lucene: does stemming of inverted index

ES: Takes care of sharding, coordinating with shards

Google Search*: uses lots of signals

- Keyword match
- Personalisation
- Page ranking algorithms (Domain authority/ Page Score, Backlinks)
- Removing Typo
- At Scaler we use ES for finding keywords in resume

Image search

- search via tags: ML models to apply tags to image, show images in captcha to train these ML models
- Not good if we need consistent data

Further reading:

1. <https://blog.insightdatascience.com/anatomy-of-an-elasticsearch-cluster-part-i-7ac9a13b05db>
2. <https://buildingvts.com/elasticsearch-architectural-overview-a35d3910e515>
3. Official Docs:
<https://www.elastic.co/guide/en/elasticsearch/reference/current/index.html>